



Akhila ✓

GenAI Trainer

**Master these 50 ML & LLM
interview questions and
land your dream AI & Tech role.**



Akhila ✓

GenAI Trainer

LLM Systems

1. How would you design a scalable AI chatbot architecture?
2. How would you reduce latency in an LLM application?
3. How would you handle rate limits for millions of users?
4. How would you cache LLM responses efficiently?
5. How would you manage conversation memory in chatbots?
6. How would you monitor hallucinations in production?
7. How would you evaluate LLM output quality automatically?
8. How would you design a multi-agent AI system?
9. How would you handle prompt injection attacks?
10. How would you deploy an open-source LLM securely?



Akhila ✓

GenAI Trainer

**We've curated a complete AI
Interview Preparation Bundle with
22 courses & 1100+ questions.**

Check out the bundle link in caption.



Akhila ✓

GenAI Trainer

RAG & Vector Databases

1. Why is chunking important in RAG systems?
2. How do you choose chunk size for documents?
3. Difference between sparse and dense retrieval
4. What causes poor retrieval quality in RAG pipelines?
5. How do embeddings work internally?
6. How would you improve semantic search accuracy?
7. Why do vector databases use similarity search?
8. Difference between cosine similarity and Euclidean distance
9. How would you evaluate retrieval performance?
10. What are rerankers and why are they used?



Akhila ✓
GenAI Trainer

MLOps & Deployment

1. How would you deploy an ML model into production?
2. What is model drift and how do you detect it?
3. Difference between CI/CD and CT pipelines in ML
4. How would you version datasets and models?
5. What metrics would you monitor in production?
6. How would you roll back a bad model deployment?
7. Difference between online and batch inference
8. How would you scale GPU inference servers?
9. What causes training-serving skew?
10. How would you optimize inference costs?



Akhila ✓

GenAI Trainer

Deep Learning

1. Why do Transformers scale better than RNNs?
2. Explain attention complexity mathematically
3. Why do larger batch sizes affect generalization?
4. Explain residual connections mathematically
5. What is gradient accumulation?
6. Why does mixed precision training improve speed?
7. What causes catastrophic forgetting?
8. Difference between fine-tuning and transfer learning
9. Why do larger models show emergent abilities?
10. How would you debug unstable training loss?



Akhila ✓

GenAI Trainer

**We've curated a complete AI
Interview Preparation Bundle with
22 courses & 1100+ questions.**

Check out the bundle link in caption.



Akhila ✓
GenAI Trainer

Python / Backend

1. How would you build a scalable FastAPI backend for AI inference?
2. Difference between synchronous and asynchronous APIs
3. How would you handle concurrent inference requests?
4. What causes memory bottlenecks in Python AI systems?
5. How would you optimize API throughput?
6. What are worker processes and thread pools?
7. Difference between multiprocessing and asyncio in AI workloads
8. How would you stream LLM responses token by token?
9. How would you design logging for AI systems?
10. How would you secure an AI API in production?



Akhila ✓
GenAI Trainer

Thanks for reading!

Save this thread and share with a friend.



Akhila ✓

GenAI Trainer

**We've curated a complete AI
Interview Preparation Bundle with
22 courses & 1100+ questions.**

Check out the bundle link in caption.